

[Components](#)[Document loaders](#)[Unstructured File](#)

# Unstructured File

This notebook covers how to use `Unstructured` package to load files of many types.

`Unstructured` currently supports loading of text files, powerpoints, html, pdfs, images, and more.

```
# # Install package
pip install "unstructured[all-docs]"
```

```
# # Install other dependencies
# # https://github.com/Unstructured-
# IO/unstructured/blob/main/docs/source/installing.rst
# !brew install libmagic
# !brew install poppler
# !brew install tesseract
# # If parsing xml / html documents:
# !brew install libxml2
# !brew install libxslt
```

```
# import nltk
# nltk.download('punkt')
```

```
from langchain.document_loaders import UnstructuredFileLoader
```

## API Reference:

- `UnstructuredFileLoader`

```
loader = UnstructuredFileLoader("./example_data/state_of_the_union.txt")
```

```
docs = loader.load()
```

```
docs[0].page_content[:400]
```

```
'Madam Speaker, Madam Vice President, our First Lady and Second Gentleman.
Members of Congress and the Cabinet. Justices of the Supreme Court. My fellow
Americans.\n\nLast year COVID-19 kept us apart. This year we are finally
together again.\n\nTonight, we meet as Democrats Republicans and Independents.
But most importantly as Americans.\n\nWith a duty to one another to the
American people to the Constit'
```

## Retain Elements

Under the hood, Unstructured creates different "elements" for different chunks of text. By default we combine those together, but you can easily keep that separation by specifying

```
mode="elements".
```

```
loader = UnstructuredFileLoader(
    "./example_data/state_of_the_union.txt", mode="elements"
)
```

```
docs = loader.load()
```

```
docs[:5]
```

```
[Document(page_content='Madam Speaker, Madam Vice President, our First
Lady and Second Gentleman. Members of Congress and the Cabinet. Justices of
the Supreme Court. My fellow Americans.', lookup_str='', metadata={'source':
'../../state_of_the_union.txt'}, lookup_index=0),
 Document(page_content='Last year COVID-19 kept us apart. This year we are
finally together again.', lookup_str='', metadata={'source':
'../../state_of_the_union.txt'}, lookup_index=0),
 Document(page_content='Tonight, we meet as Democrats Republicans and
Independents. But most importantly as Americans.', lookup_str='', metadata=
{'source': '../../state_of_the_union.txt'}, lookup_index=0),
 Document(page_content='With a duty to one another to the American people
to the Constitution.', lookup_str='', metadata={'source':
'../../state_of_the_union.txt'}, lookup_index=0),
 Document(page_content='And with an unwavering resolve that freedom will
always triumph over tyranny.', lookup_str='', metadata={'source':
'../../state_of_the_union.txt'}, lookup_index=0)]
```

# Define a Partitioning Strategy

Unstructured document loader allow users to pass in a `strategy` parameter that lets `unstructured` know how to partition the document. Currently supported strategies are `"hi_res"` (the default) and `"fast"`. Hi res partitioning strategies are more accurate, but take longer to process. Fast strategies partition the document more quickly, but trade-off accuracy. Not all document types have separate hi res and fast partitioning strategies. For those document types, the `strategy` kwarg is ignored. In some cases, the high res strategy will fallback to fast if there is a dependency missing (i.e. a model for document partitioning). You can see how to apply a strategy to an `UnstructuredFileLoader` below.

```
from langchain.document_loaders import UnstructuredFileLoader
```

## API Reference:

- [UnstructuredFileLoader](#)

```
loader = UnstructuredFileLoader(  
    "layout-parser-paper-fast.pdf", strategy="fast", mode="elements"  
)
```

```
docs = loader.load()
```

```
docs[:5]
```

```
[Document(page_content='1', lookup_str='', metadata={'source': 'layout-  
parser-paper-fast.pdf', 'filename': 'layout-parser-paper-fast.pdf',  
'page_number': 1, 'category': 'UncategorizedText'}, lookup_index=0),  
 Document(page_content='2', lookup_str='', metadata={'source': 'layout-  
parser-paper-fast.pdf', 'filename': 'layout-parser-paper-fast.pdf',  
'page_number': 1, 'category': 'UncategorizedText'}, lookup_index=0),  
 Document(page_content='0', lookup_str='', metadata={'source': 'layout-  
parser-paper-fast.pdf', 'filename': 'layout-parser-paper-fast.pdf',  
'page_number': 1, 'category': 'UncategorizedText'}, lookup_index=0),  
 Document(page_content='2', lookup_str='', metadata={'source': 'layout-  
parser-paper-fast.pdf', 'filename': 'layout-parser-paper-fast.pdf',  
'page_number': 1, 'category': 'UncategorizedText'}, lookup_index=0),  
 Document(page_content='n', lookup_str='', metadata={'source': 'layout-
```

```
parser-paper-fast.pdf', 'filename': 'layout-parser-paper-fast.pdf',
'page_number': 1, 'category': 'Title'}, lookup_index=0)]
```

## PDF Example

Processing PDF documents works exactly the same way. Unstructured detects the file type and extracts the same types of elements. Modes of operation are

- `single` all the text from all elements are combined into one (default)
- `elements` maintain individual elements
- `paged` texts from each page are only combined

```
wget https://raw.githubusercontent.com/Unstructured-
IO/unstructured/main/example-docs/layout-parser-paper.pdf -P "../.."
```

```
loader = UnstructuredFileLoader(
    "../example_data/layout-parser-paper.pdf", mode="elements"
)
```

```
docs = loader.load()
```

```
docs[:5]
```

```
[Document(page_content='LayoutParser : A Unified Toolkit for Deep Learning
Based Document Image Analysis', lookup_str='', metadata={'source':
'../..../layout-parser-paper.pdf'}, lookup_index=0),
 Document(page_content='Zejiang Shen 1 ( (ea)\n ), Ruochen Zhang 2 ,
Melissa Dell 3 , Benjamin Charles Germain Lee 4 , Jacob Carlson 3 , and
Weining Li 5', lookup_str='', metadata={'source': '../..../layout-parser-
paper.pdf'}, lookup_index=0),
 Document(page_content='Allen Institute for AI shannons@allenai.org',
lookup_str='', metadata={'source': '../..../layout-parser-paper.pdf'},
lookup_index=0),
 Document(page_content='Brown University ruochen zhang@brown.edu',
lookup_str='', metadata={'source': '../..../layout-parser-paper.pdf'},
lookup_index=0),
 Document(page_content='Harvard University { melissadell,jacob carlson }
```

```
@fas.harvard.edu', lookup_str='', metadata={'source': '.././layout-parser-paper.pdf'}, lookup_index=0)]
```

If you need to post process the `unstructured` elements after extraction, you can pass in a list of `str` -> `str` functions to the `post_processors` kwarg when you instantiate the `UnstructuredFileLoader`. This applies to other Unstructured loaders as well. Below is an example.

```
from langchain.document_loaders import UnstructuredFileLoader
from unstructured.cleaners.core import clean_extra_whitespace
```

### API Reference:

- [UnstructuredFileLoader](#)

```
loader = UnstructuredFileLoader(
    "../example_data/layout-parser-paper.pdf",
    mode="elements",
    post_processors=[clean_extra_whitespace],
)
```

```
docs = loader.load()
```

```
docs[:5]
```

```
[Document(page_content='LayoutParser: A Unified Toolkit for Deep Learning
Based Document Image Analysis', metadata={'source': '../example_data/layout-
parser-paper.pdf', 'coordinates': {'points': ((157.62199999999999,
114.23496279999995), (157.62199999999999, 146.5141628), (457.7358962799999,
146.5141628), (457.7358962799999, 114.23496279999995)), 'system':
'PixelSpace', 'layout_width': 612, 'layout_height': 792}, 'filename': 'layout-
parser-paper.pdf', 'file_directory': '../example_data', 'filetype':
'application/pdf', 'page_number': 1, 'category': 'Title'}),
```

```
Document(page_content='Zejiang Shen1 ((cid:0)), Ruochen Zhang2, Melissa
Dell3, Benjamin Charles Germain Lee4, Jacob Carlson3, and Weining Li5',
metadata={'source': '../example_data/layout-parser-paper.pdf', 'coordinates':
{'points': ((134.809, 168.64029940800003), (134.809, 192.2517444),
(480.5464199080001, 192.2517444), (480.5464199080001, 168.64029940800003)),
'system': 'PixelSpace', 'layout_width': 612, 'layout_height': 792},
'filename': 'layout-parser-paper.pdf', 'file_directory': '../example_data',
```

```
'filetype': 'application/pdf', 'page_number': 1, 'category':
'UncategorizedText'}},
    Document(page_content='1 Allen Institute for AI shannons@allenai.org 2
Brown University ruochen zhang@brown.edu 3 Harvard University
{melissadell,jacob carlson}@fas.harvard.edu 4 University of Washington
bcgl@cs.washington.edu 5 University of Waterloo w422li@uwaterloo.ca',
metadata={'source': './example_data/layout-parser-paper.pdf', 'coordinates':
{'points': ((207.23000000000002, 202.57205439999996), (207.23000000000002,
311.8195408), (408.12676, 311.8195408), (408.12676, 202.57205439999996)),
'system': 'PixelSpace', 'layout_width': 612, 'layout_height': 792},
'filename': 'layout-parser-paper.pdf', 'file_directory': './example_data',
'filetype': 'application/pdf', 'page_number': 1, 'category':
'UncategorizedText'}},
    Document(page_content='1 2 0 2', metadata={'source':
'./example_data/layout-parser-paper.pdf', 'coordinates': {'points': ((16.34,
213.36), (16.34, 253.36), (36.34, 253.36), (36.34, 213.36)), 'system':
'PixelSpace', 'layout_width': 612, 'layout_height': 792}, 'filename': 'layout-
parser-paper.pdf', 'file_directory': './example_data', 'filetype':
'application/pdf', 'page_number': 1, 'category': 'UncategorizedText'}},
    Document(page_content='n u J', metadata={'source':
'./example_data/layout-parser-paper.pdf', 'coordinates': {'points': ((16.34,
258.36), (16.34, 286.14), (36.34, 286.14), (36.34, 258.36)), 'system':
'PixelSpace', 'layout_width': 612, 'layout_height': 792}, 'filename': 'layout-
parser-paper.pdf', 'file_directory': './example_data', 'filetype':
'application/pdf', 'page_number': 1, 'category': 'Title'}}])
```

## Unstructured API

If you want to get up and running with less set up, you can simply run `pip install unstructured` and use `UnstructuredAPIFileLoader` or `UnstructuredAPIFileIOLoader`. That will process your document using the hosted Unstructured API. You can generate a free Unstructured API key [here](#). The [Unstructured documentation](#) page will have instructions on how to generate an API key once they're available. Check out the instructions [here](#) if you'd like to self-host the Unstructured API or run it locally.

```
from langchain.document_loaders import UnstructuredAPIFileLoader
```

### API Reference:

- [UnstructuredAPIFileLoader](#)

```
filenames = ["example_data/fake.docx", "example_data/fake-email.eml"]
```

```
loader = UnstructuredAPIFileLoader(  
    file_path=filenames[0],  
    api_key="FAKE_API_KEY",  
)
```

```
docs = loader.load()  
docs[0]
```

```
Document(page_content='Lorem ipsum dolor sit amet.', metadata={'source':  
'example_data/fake.docx'})
```

You can also batch multiple files through the Unstructured API in a single API using

`UnstructuredAPIFileLoader`.

```
loader = UnstructuredAPIFileLoader(  
    file_path=filenames,  
    api_key="FAKE_API_KEY",  
)
```

```
docs = loader.load()  
docs[0]
```

```
Document(page_content='Lorem ipsum dolor sit amet.\n\nThis is a test email  
to use for unit tests.\n\nImportant points:\n\nRoses are red\n\nViolets are  
blue', metadata={'source': ['example_data/fake.docx', 'example_data/fake-  
email.eml']})
```